


> stride = 1

1 load
100% efficiency

> All 32 threads land in
the same cache line,
one fetch covers all elements
Memory access is coalesced



2 2 loads, 50% efficiency

3 3 loads, 33% efficiency

> stride = 4

4 loads
25% efficiency

> Threads span 4 cache
lines, so only 25% of
fetched data is actually used
This is noncoalesced.

